

Analytics Systems Technology



ALY6140, WINTER 2022

Capstone Project Draft

Capstone Group: 09

Submitted by: Abhinav Jain

Submitted To: Prof. Richard Zhi

Date: 03/31/2022

Football Match Probability Prediction

Introduction:

Predicting match outcomes is the most difficult undertaking, and it always delights football fans to try to anticipate the conclusion of each match. In the discipline of data science, football takes center stage. We will create models to forecast match results using event recognition and team analysis in this paper. The outcome of a match between two teams is mostly determined by their current form, but the teams' current form may be determined by looking at their recent sequence results against other teams. As a result, the match's probability might differ between two teams.

William McGregor, a Perthshire lord, and Aston Villa F.C. director was the driving force behind meetings in London and Manchester in 1888 with 12 football teams in the hopes of organizing a league. These 12 clubs would eventually become the Football League's 12 founding members.

To forecast the outcome of a football game, the project dataset was acquired from Kaggle.com's "Football Match Probability Prediction." The information was gathered from 150000 past international football matches between 2019 and 2021 by the community prediction competition team. The information was gathered from over 9500 teams and 860 leagues. We'll choose 10-15 variables from this dataset, which comprises 189 variables and 72711 observations, depending on the project's needs, in order to forecast and show the results using multiple regression models.

Football is constantly in high demand since there are sports fanatics all around the world. Fans of football are nervous and interested in the game's uncertain outcomes, which adds to the thrills and tension. Football is the most fun activity for fans to divert their attention away from their daily concerns. Leagues will allow us to deduce patterns from the previous success and failure of matches, as well as another variable in this project team's data, to predict the outcome. It's based on a number of factors, including target, home_team_name, away_team_name, player id, coach id, league id, match date, ratings and other various factors.

Goals:

To assess the chance of a future event, data analytics, statistical algorithms, and machine learning approaches are used to compare data to previous events. Instead of only knowing what has happened, the goal is to make better predictions of what will happen in the future.

Predictive techniques are used to analyze the team's performance during the match. In this study, different models will be employed to estimate the results of a football match, including Random Forest Classifier, K Nearest Neighbor, and Logistics Regression.

The purpose of a data scientist is to collect data, do analysis, interpret the data in a meaningful way, and apply prediction models to the data.

Methods:

I: Predict the matches by getting historical data of the team with the home team coach and opponent team coach.

II. Predict the accuracy by knowing the history match dates of the home and opponent team

III. Predict the motivation of the team by evaluating the team history rating.

Questions to Investigate :

1. What was the team performance history at home play?
2. What was the rating of the opponent team?
3. In the past, when did a match help you?
4. Which coach had a better track record in previous leagues?
5. What were the results of a team's prior leagues?

Exploratory Data Analysis

The process of exploratory data analysis helps in analyzing the insights from the dataset which provide the statistical analysis and graphical representation of the data. In this dataset for the prediction of a football match from the huge dataset, a selected few variables were to apply the predictive models. Initially, after importing the dataset in a jupyter notebook, installing various modules which will require analyzing the dataset helped in the visualization of the statistics.

Description analysis:

In this task, the shape of the dataset consists of 110938 rows and 190 columns, whereas id is an integer, target, home team name, away team name, match date, away team history are float datatype. Selection of the variable is the most difficult task to undertake with the predictive models in mind.

- Shape of the dataset:

```
df.shape
```

```
(110938, 190)
```

- Datatypes of each variable

```
df.dtypes
```

```
id                int64
target            object
home_team_name    object
away_team_name    object
match_date        object
...
away_team_history_league_id_6  float64
away_team_history_league_id_7  float64
away_team_history_league_id_8  float64
away_team_history_league_id_9  float64
away_team_history_league_id_10 float64
Length: 190, dtype: object
```

Some common features used from the dataset:

1. Id: Match identity of the data available for the team
2. Target: this feature of the dataset will compare the home team and the away team
3. Home_team_name: team name which is called Home
4. Away_team_name: team name which defines away
5. Match_date: define the time zone of the match
6. League_name: it explains the name of the league
7. League_id: for two team league names can be found identical
8. Is_cup : the match played for the cup
9. Home_team_coach_id: it defines the id of the coach in Hometeam
10. Away_team_coach_id: it defines the id of the coach in the away team

There are more features that we have to use for the comparison to implement the predictive model in the dataset. The historical data are taken from the more than 150000 historical world football matches from 2019-2021 which define our goal in forecasting the outcome of the match by implementing the various models.

Data Extraction

After importing the data from the dataset. These features include historical data of the home team and the away team

```
df.head()
```

	id	target	home_team_name	away_team_name	match_date	league_name	league_id	is_cup	home_team_coach_id	away_team_coach_id	...	away_te
0	11906497	away	Newell's Old Boys	River Plate	2019-12-01 00:45:00	Superliga	636	False	468196.00000	468200.00000	...	
1	11984383	home	Real Estelí	Deportivo Las Sabanias	2019-12-01 01:00:00	Primera Division	752	False	516788.00000	22169161.00000	...	
2	11983301	draw	UPNFM	Marathón	2019-12-01 01:00:00	Liga Nacional	734	False	2510608.00000	456313.00000	...	
3	11983471	away	León	Morelia	2019-12-01 01:00:00	Liga MX	743	False	1552508.00000	465797.00000	...	
4	11883005	home	Cobán Imperial	Iztapa	2019-12-01 01:00:00	Liga Nacional	705	False	429958.00000	426870.00000	...	

5 rows × 190 columns

```
df.shape
```

(110938, 190)

Data Cleanup: In this task after analyzing the dataset, we drop many variables from the dataset and took variables as per our requirements which will help in predicting the accuracy of the model are given below:

- 'home_team_history_is_play_home_1'
- 'home_team_history_opponent_goal_1'
- 'home_team_history_rating_1'
- 'home_team_history_coach_1'
- 'home_team_history_league_id_1'
- 'away_team_history_is_play_home_1'
- 'away_team_history_goal_1'
- 'away_team_history_opponent_goal_1'
- 'away_team_history_rating_1'
- 'away_team_history_opponent_rating_1'
- 'away_team_history_coach_1'

```
data= df[[
    'target', 'home_team_name', 'away_team_name', 'league_name',
    'home_team_coach_id', 'away_team_coach_id', 'home_team_history_is_play_home_1',
    'home_team_history_opponent_goal_1', 'home_team_history_rating_1', 'home_team_history_coach_1',
    'home_team_history_league_id_1', 'away_team_history_is_play_home_1',
    'away_team_history_goal_1', 'away_team_history_opponent_goal_1', 'away_team_history_rating_1',
    'away_team_history_opponent_rating_1', 'away_team_history_coach_1'
]].copy()

data.head()
```

	target	home_team_name	away_team_name	league_name	home_team_coach_id	away_team_coach_id	home_
0	away	Newell's Old Boys	River Plate	Superliga	468196.00000	468200.00000	
1	home	Real Estelí	Deportivo Las Sabanas	Primera Division	516788.00000	22169161.00000	
2	draw	UPNFM	Marathón	Liga Nacional	2510608.00000	456313.00000	
3	away	León	Morelia	Liga MX	1552508.00000	465797.00000	
4	home	Cobán Imperial	Iztapa	Liga Nacional	429958.00000	426870.00000	

```
data.shape
```

```
(110938, 17)
```


```
df.dtypes
```

```
id                int64
target            object
home_team_name    object
away_team_name    object
match_date        object
...
away_team_history_league_id_6    float64
away_team_history_league_id_7    float64
away_team_history_league_id_8    float64
away_team_history_league_id_9    float64
away_team_history_league_id_10   float64
Length: 190, dtype: object
```

Data Analysis and Interpretation:


1. In this task analyses the comparison of the home team with the name and target of the team. The below bar graph is shown to describe the home team name with counts with comparison away, home, draw as we are targeting for the prediction of the match.

```
plt.figure(figsize=(10,10))
sns.countplot(x="home_team_name",hue="target",data=df,order=df.home_team_name.value_counts().iloc[:5].index)
plt.show()
```



2. In this task analyses the comparison of the away team with name and target of the team The below bar graph is shown to describe the away team name with counts with comparison away, home, draw as we are targeting for the prediction of the match.

```
plt.figure(figsize=(10,10))
sns.countplot(x="away_team_name",hue="target",data=df,order=df.away_team_name.value_counts().iloc[:5].index)
plt.show()
```



3. In this task collect the data of the team history at home team and away team

```
df.groupby('home_team_name')[play_home_features].mean()
```


	home_team_history_is_play_home_1	home_team_history_is_play_home_2	home_team_history_is_play_home_3
home_team_name			
07 Vestur	0.53846	0.38462	0.53846
1. FC M'gladbach	0.00000	1.00000	0.00000
1. FC Merseburg	0.00000	0.66667	0.50000
1. Maj Ruma	0.20000	0.50000	0.50000
12 de Octubre	0.32000	0.64000	0.44000
...
Žilina	0.34483	0.48276	0.55172
Žilina II	0.30000	0.65000	0.50000
Žilina U19	0.14286	0.71429	0.28571
Župa	0.00000	0.50000	0.00000
Žďár nad Sázavou	0.25000	0.50000	0.25000

9813 rows × 20 columns

```
df.groupby('home_team_name')[opponent_rating_features].mean()
```


	home_team_history_opponent_rating_1	home_team_history_opponent_rating_2	home_team_history_opponent_r
home_team_name			
07 Vestur	8.61168	8.75077	
1. FC M'gladbach	11.37313	9.09024	11
1. FC Merseburg	9.23196	7.66353	
1. Maj Ruma	7.18673	5.48778	
12 de Octubre	7.52886	7.21602	
...	
Žilina	6.19983	6.05111	
Žilina II	8.96911	7.35407	
Žilina U19	7.32640	7.03376	
Župa	8.55600	6.40655	11
Žďár nad Sázavou	11.56933	12.10130	11

9813 rows × 20 columns



4. In this task analyze the track record of the coach in the previous matches with the comparison between away, home teams, and draw.

```
plt.figure(figsize=(10,10))
sns.countplot(x="away_team_coach_id",hue="target",data=df,order=df.away_team_coach_id.value_counts().index)
plt.show()
```



5. In this task analyze the track record of the history cup of the away team

```
df.groupby('away_team_name')[away_is_cup_features].mean()
```

away_team_name	away_team_history_is_cup_1	away_team_history_is_cup_2	away_team_history_is_cup_3	away_team_history_is
07 Vestur	0.00000	0.00000	0.11111	
1. FC M'gladbach	0.00000	0.00000	0.00000	
1. FC Merseburg	0.00000	0.00000	0.00000	
1. Maj Ruma	0.00000	0.00000	0.00000	
12 de Octubre	0.11538	0.07692	0.03846	
...	
Žilina	0.10714	0.10714	0.21429	
Žilina II	0.00000	0.00000	0.00000	
Žilina U19	0.00000	0.00000	0.00000	
Župa	0.00000	0.00000	0.00000	
Žďár nad Sázavou	0.00000	0.00000	0.00000	

9892 rows × 10 columns

Predictive Models: In the present job market, the future of predictive modeling in statistics is the key emphasis. We may forecast the answer via predictive modeling by adopting modeling solutions based on historical and recent data. Data is obtained and created using predictions, and then confirmed using the most up-to-date information available once the dataset has been trained.


1. **Logistic Regression Model** When the variable is categorical, what type of regression analysis should be used? (binary). The logistic regression, like other regression studies, is a predictive analysis. To describe data and explain the connection between one dependent binary variable and one or more nominal, ordinal, interval, or ratio-level independent variables, logistic regression is utilized.
 - Easier to Implement
 - Efficient to train
 - Fast while classifying unknown records
 - Interpret the Model coefficient as indicators

Logistics Regression helps in predictive analysis to evaluate the relationship between the nominal ordinal of the interval, by checking the ratio-level independent variables, or one dependent binary variable. The regression studies suggest that the cross-validation is 1.0462 with 46% of accuracy.

Accuracy score: 0.4908959798089057

	precision	recall	f1-score	support
away	0.48	0.44	0.46	7035
draw	0.42	0.00	0.01	5530
home	0.50	0.81	0.62	9623
accuracy			0.49	22188
macro avg	0.47	0.42	0.36	22188
weighted avg	0.47	0.49	0.41	22188

```
# Confusion matrix
cm = confusion_matrix(y_test, y_pred, labels=pipeline.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=pipeline.classes_)
fig, ax = plt.subplots(figsize=(15,7))
disp.plot(ax=ax)
plt.show()
```



2. **Random Forest Classifier:** is a classifier that combines a number of decision trees on different subsets of a dataset and averages the results to increase the dataset's predicted accuracy." Rather than depending on a single decision tree, the random forest collects predictions from each tree and makes decisions based on them.


- Works well with large dimensional data
- Working with a subset
- Fast to train than the decision tree
- Easily work with hundreds of features
- Low correlation is the key

Random Forest classifier is based on supervised learning and provides the solution for regression and creating multiple decision trees by evaluating the average prediction of that tree Here, the cross-validation is 1.0462 with 46% of accuracy.

Accuracy score: 0.46705426356589147

	precision	recall	f1-score	support
away	0.57	0.17	0.26	7035
draw	0.00	0.00	0.00	5530
home	0.46	0.95	0.62	9623
accuracy			0.47	22188
macro avg	0.34	0.37	0.29	22188
weighted avg	0.38	0.47	0.35	22188

```
# Confusion matrix
cm = confusion_matrix(y_test, y_pred, labels=pipeline.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=pipeline.classes_)
fig, ax = plt.subplots(figsize=(15,7))
disp.plot(ax=ax)
plt.show()
```



3. **KNN-neighbors:** The supervised learning technique K-nearest neighbors (KNN) is used for both regression and classification. By computing the distance between the test data and all of the training points, KNN tries to predict the proper class for the test data. Then choose the K number of points that are the most similar to the test data.


- Highly accurate Predictions
- Solves both classification and regression problem statement
- KNN algorithm for multiclass classification
- Recommendation Systems

KNN regression model is defined by the analyst by measuring the size of the neighborhood to check the validation of the model. Here, the cross-validation is 9.79.17 with 39% of accuracy

Accuracy score: 0.3992248062015504

	precision	recall	f1-score	support
away	0.36	0.49	0.42	7035
draw	0.26	0.17	0.20	5530
home	0.49	0.46	0.48	9623
accuracy			0.40	22188
macro avg	0.37	0.37	0.37	22188
weighted avg	0.39	0.40	0.39	22188

```
# Confusion matrix
cm = confusion_matrix(y_test, y_pred, labels=pipeline.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=pipeline.classes_)
fig, ax = plt.subplots(figsize=(15,7))
disp.plot(ax=ax)
plt.show()
```



Interpretive and Conclusions

Logistic Regression confusion matrix provides the data predictive label which includes home team, away that is opponent team, and draw team which has matched with home and away has 1807 matches, 9 of home and draw, home and home has 7807, draw and away has 1550, draw and draw 14 matches, draw at home has 3966, away and away 3071 matches, away and draw has 10 matches, away and home has 3954 matches which leads to the 49% of accuracy. However, the cross-validation log loss is 1.0208.

Random classifier confusion matrix provides the data predictive label which includes home team, away that is opponent team, and draw team which has matched with home and away has 435 matches, 0 of home and draw, home and home has 9188, draw and away has 437, draw and draw 0 matches, draw at home has 5093, away and away 1175 matches, away and draw has 0 matches, away and home has 5860 matches which leads to the 46% of accuracy. In this case, cross-validation log loss is 0.0462

KNN confusion matrix provides the data predictive label which includes home team, away that is opponent team, and draw team which has matched with home and away has 3660 matches, 1507 of home and draw, home and home has 4456, draw and away has 2407, draw and draw 933 matches, draw at home has 2190, away and away 3469 matches, away and draw has 1172 matches, away and home has 2394 matches which lead to the 39% of accuracy. In this model, we can see the cross-validation is 9.7917.

In conclusion, we experimented with a variety of Machine Learning models and features to improve our prediction error as much as possible. The predicting outcome will help determine the instance outcome of the psychology towards the match and create the opportunity. Many organizations are using these forecasting methods to build for the outperformance of the results in the field of data science.

We removed numerous factors from the dataset after evaluating it and replaced them with variables that will help us anticipate the model's accuracy. After that compares the home team's name and target to the team's name and target. The bar graph below depicts the home team name with counts in relation to away, home, and draw since we are aiming for a match prediction. Compares the away side's name and goal to that of the home team. The bar graph below depicts the away team's name with counts and a comparison of away, home, and draw, as we are aiming for a match prediction.

Logistics Regression helps in the evaluation of the link between the nominal ordinal of the interval and the ratio-level independent variables, or one dependent binary variable, in predictive analysis. According to regression research, the cross-validation is 1.0462 with a 46 percent accuracy.

The Random Forest classifier is based on supervised learning and provides a solution for regression and multiple decision tree creation by evaluating the tree's average prediction. The cross-validation, in this case, is 1.0462 with a 46 percent accuracy.

K-Nearest Neighbor The analyst defines the KNN regression model by measuring the size of the neighborhood to ensure the model's validity. The cross-validation score is 9.79.17, with a precision of 39%.

We built a model training and testing pipeline using our various model components to quickly and easily change and test new assumptions. We compared our forecasts to benchmark methodologies to better understand the predictive performance of our models. We discovered that this model may be improved to obtain data on the matches and the coach id, which appears to be an outlier. We may have predicted the models with less accuracy due to a lack of data. We need to anticipate the train and test the model in such a way that we can reach that accuracy to develop the model to the level of 90-95 percent accuracy, and we need to provide a better outcome after running the model numerous times.

=
=====

References:

[1]Football Match Probability Prediction

<https://www.kaggle.com/c/football-match-probability-prediction/data?select=test.csv>

[2]learn

<https://scikit-learn.org/stable/>

[3]pandas

<https://pandas.pydata.org/>

[4]seaborn: statistical data visualization — seaborn 0.11.2 ...

seaborn: statistical data visualization. ¶. Seaborn is a Python data visualization library based on matplotlib. It provides a high-level interface for drawing attractive and informative statistical graphics. For a brief introduction to the ideas behind the library, you can read the introductory notes or the paper.

<https://seaborn.pydata.org>

[5]sklearn.ensemble.RandomForestClassifier

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>

[6]sklearn.neighbors.KNeighborsClassifier

<https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>

[7]sklearn.model_selection.train_test_split

https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html

[8]sklearn.metrics.accuracy_score

https://scikit-learn.org/stable/modules/generated/sklearn.metrics.accuracy_score.html

[9]sklearn.metrics.confusion_matrix

https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html

[10]sklearn.metrics.classification_report

https://scikit-learn.org/stable/modules/generated/sklearn.metrics.classification_report.html

[11]warnings - Warning control¶

<https://docs.python.org/3/library/warnings.html>

[12]sklearn.linear_model.LogisticRegression

https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html

Appendix:

```
1  #!/usr/bin/env python
2  # coding: utf-8
3
4  Run Cell | Run Below | Debug Cell
5  # In[1]:
6
7  #import modules
8  from sklearn import metrics, model_selection
9  from sklearn.pipeline import Pipeline
10 from sklearn.linear_model import LogisticRegression
11 from sklearn.impute import SimpleImputer
12 import pandas as pd
13 import numpy as np
14 import matplotlib.pyplot as plt
15 import seaborn as sns
16 from sklearn.ensemble import RandomForestClassifier
17 from sklearn.neighbors import KNeighborsClassifier
18 from sklearn.model_selection import train_test_split
19 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, ConfusionMatrixDisplay
20 pd.options.display.float_format = '{:.5f}'.format
21
22 get_ipython().run_line_magic('matplotlib', 'inline')
23 import warnings
24 warnings.filterwarnings('ignore')
25
26 Run Cell | Run Above | Debug Cell
27 # In[2]:
28
29
30 df = pd.read_csv("train.csv")
31
```

```
Run Cell | Run Above | Debug Cell
33 # In[3]:
34
35
36 df.head()
37
38
Run Cell | Run Above | Debug Cell
39 # In[4]:
40
41
42 df.shape
43
44
Run Cell | Run Above | Debug Cell
45 # In[5]:
46
47
48 df.columns
49
50
Run Cell | Run Above | Debug Cell
51 # In[6]:
52
53
54 df['target']
55
56
Run Cell | Run Above | Debug Cell
57 # In[7]:
58
59
60 df['home_team_history_rating_5'].value_counts()
61
62
```

```

Run Cell | Run Above | Debug Cell
63 # In[8]:
64
65
66 df['league_id'].value_counts()
67
68
Run Cell | Run Above | Debug Cell
69 # In[9]:
70
71
72 #keep only rating features
73 rating_features = [x for x in df if 'rating' in x]
74 rating_features
75
76
Run Cell | Run Above | Debug Cell
77 # In[ ]:
78
79
80
81
82
Run Cell | Run Above | Debug Cell
83 # In[10]:
84
85
86 data= df[[
87     'target', 'home_team_name', 'away_team_name', 'league_name',
88     'home_team_coach_id', 'away_team_coach_id', 'home_team_history_is_play_home_1',
89     'home_team_history_opponent_goal_1', 'home_team_history_rating_1', 'home_team_history_coach_1',
90     'home_team_history_league_id_1', 'away_team_history_is_play_home_1',
91     'away_team_history_goal_1', 'away_team_history_opponent_goal_1', 'away_team_history_rating_1',
92     'away_team_history_opponent_rating_1', 'away_team_history_coach_1'
93 ]].copy()
94

```

```

93     ]].copy()
94
95     data.head()
96
97     Run Cell | Run Above | Debug Cell
98     # In[11]:
99
100
101     data.shape
102
103
104     Run Cell | Run Above | Debug Cell
105     # In[12]:
106
107     df.dtypes
108
109
110     Run Cell | Run Above | Debug Cell
111     # In[13]:
112
113     df['away_team_history_coach_1'].value_counts()
114
115
116     Run Cell | Run Above | Debug Cell
117     # In[14]:
118
119     df['home_team_history_is_play_home_1'].value_counts()
120

```

```

Run Cell | Run Above | Debug Cell
122 # In[15]:
123
124
125 plt.figure(figsize=(10,10))
126 sns.countplot(x="home_team_name",hue="target",data=df,order=df.home_team_name.value_counts().iloc[:5].index)
127 plt.show()
128
129
Run Cell | Run Above | Debug Cell
130 # In[16]:
131
132
133 plt.figure(figsize=(10,10))
134 sns.countplot(x="away_team_name",hue="target",data=df,order=df.away_team_name.value_counts().iloc[:5].index)
135 plt.show()
136
137
138 # ## Team history at home
139
Run Cell | Run Above | Debug Cell
140 # In[17]:
141
142
143 #keep only play home history features
144 play_home_features = [x for x in df if 'play_home' in x]
145 play_home_features
146
147
Run Cell | Run Above | Debug Cell
148 # In[18]:
149
150
151 df.groupby('home_team_name')[play_home_features].mean()

```

```

150
151 df.groupby('home_team_name')[play_home_features].mean()
152
153
154 # ### Rating of the opponent team
155
    Run Cell | Run Above | Debug Cell
156 # In[19]:
157
158
159 #keep only opponent rating features
160 opponent_rating_features = [x for x in df if 'opponent_rating' in x]
161 opponent_rating_features
162
163
    Run Cell | Run Above | Debug Cell
164 # In[20]:
165
166
167 df.groupby('home_team_name')[opponent_rating_features].mean()
168
169
170 # ### In the past, when did a match help
171
    Run Cell | Run Above | Debug Cell
172 # In[21]:
173
174
175 #keep only home team history mach date features
176 home_date_features = [x for x in df if 'home_team_history_match_date' in x]
177 home_date_features

```

```

183 df['target'].value_counts()
184
185
186 Run Cell | Run Above | Debug Cell
187 # In[44]:
188
189 home_wins = df[df['target'] == "home"]
190
191
192 Run Cell | Run Above | Debug Cell
193 # In[47]:
194
195 home_wins[home_date_features]
196
197
198 # ## Which coach had a better track record
199
200 Run Cell | Run Above | Debug Cell
201 # In[23]:
202
203 away_coach_features = [x for x in df if 'away_team_history_coach' in x]
204 away_coach_features
205
206
207 Run Cell | Run Above | Debug Cell
208 # In[24]:
209
210 plt.figure(figsize=(10,10))
211 sns.countplot(x="away_team_coach_id",hue="target",data=df,order=df.away_team_coach_id.value_counts().iloc[:5].index)
212 plt.show()

```



```

215 # ## Teams' prior league results
216
217 Run Cell | Run Above | Debug Cell
218 # In[25]:
219
220 #keep only away history cup
221 away_is_cup_features = [x for x in df if 'away_team_history_is_cup' in x]
222 away_is_cup_features
223
224
225 Run Cell | Run Above | Debug Cell
226 # In[26]:
227
228 df.groupby('away_team_name')[away_is_cup_features].mean()
229
230
231 # ## Modelling
232
233 Run Cell | Run Above | Debug Cell
234 # In[27]:
235
236 #Make X and y
237 X = df[rating_features]
238 y = df['target']
239
240
241 Run Cell | Run Above | Debug Cell
242 # In[28]:
243
244 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

```

```

247 # ### Logistic Regression
248
249 # In[29]:
250
251
252 #build a simple pipeline that fill nans values with the mean and fit a logistic regression classifier without regulariation.
253 pipeline = Pipeline(steps=[
254     ('imputer', SimpleImputer(strategy='mean')),
255     ("classifier", LogisticRegression(solver='sag',C=1e5))
256 ])
257
258
259 # In[30]:
260
261
262 #evaluate the model log loss using cross validation
263 cv_scores = model_selection.cross_val_score(pipeline, X_train, y_train, cv=10,scoring='neg_log_loss')
264 print(f'The cross validation log loss is {-cv_scores.mean().round(4)}')
265
266
267 # In[31]:
268
269
270 #fit the model
271 pipeline.fit(X_train, y_train)
272
273 # Making predictions
274 y_pred = pipeline.predict(X_test)
275
276 # Measuring the accuracy of the model
277 print(f'Accuracy score: {accuracy_score(y_test, y_pred)}')
278 print('\n')
279 print(f'{classification_report(y_test, y_pred)}')
280
281 # In[32]:
282
283
284
285 # Confusion matrix
286 cm = confusion_matrix(y_test, y_pred, labels=pipeline.classes_)
287 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=pipeline.classes_)
288 fig, ax = plt.subplots(figsize=(15,7))
289 disp.plot(ax=ax)
290 plt.show()
291
292
293 # ## Random Classifier
294
295 # In[33]:
296
297
298 #build a simple pipeline that fill nans values with the mean and fit a logistic regression classifier without regulariation.
299 pipeline = Pipeline(steps=[
300     ('imputer', SimpleImputer(strategy='mean')),
301     ("classifier", RandomForestClassifier(max_depth = 3, random_state = 0))
302 ])
303
304
305 # In[34]:
306
307
308 #evaluate the model log loss using cross validation
309 cv_scores = model_selection.cross_val_score(pipeline, X_train, y_train, cv=10,scoring='neg_log_loss')
310 print(f'The cross validation log loss is {-cv_scores.mean().round(4)}')

```

```

322 #FIT THE MODEL
323 pipeline.fit(X_train, y_train)
324
325 # Making predictions
326 y_pred = pipeline.predict(X_test)
327
328 # Measuring the accuracy of the model
329 print(f'Accuracy score: {accuracy_score(y_test, y_pred)}')
330 print('\n')
331 print(f'{classification_report(y_test, y_pred)}')
332
333
Run Cell | Run Above | Debug Cell
334 # In[36]:
335
336
337 # Confusion matrix
338 cm = confusion_matrix(y_test, y_pred, labels=pipeline.classes_)
339 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=pipeline.classes_)
340 fig, ax = plt.subplots(figsize=(15,7))
341 disp.plot(ax=ax)
342 plt.show()
343
344
345 # ## kNEIGHBORS
346
Run Cell | Run Above | Debug Cell
347 # In[37]:
348
349
350 #build a simple pipeline that fill nans values with the mean and fit a logistic regression classifier without regulariation.
351 pipeline = Pipeline(steps=[
352     ('imputer', SimpleImputer(strategy='mean')),
353     ('classifier", KNeighborsClassifier(n_neighbors=3))
354 ])

```

```

Run Cell | Run Above | Debug Cell
357 # In[38]:
358
359
360 #evaluate the model log loss using cross validation
361 cv_scores = model_selection.cross_val_score(pipeline, X_train, y_train, cv=10,scoring='neg_log_loss')
362 print(f'The cross validation log loss is {-cv_scores.mean().round(4)}')
363
364
Run Cell | Run Above | Debug Cell
365 # In[39]:
366
367
368 #FIT THE MODEL
369 pipeline.fit(X_train, y_train)
370
371 # Making predictions
372 y_pred = pipeline.predict(X_test)
373
374 # Measuring the accuracy of the model
375 print(f'Accuracy score: {accuracy_score(y_test, y_pred)}')
376 print('\n')
377 print(f'{classification_report(y_test, y_pred)}')
378
379
Run Cell | Run Above | Debug Cell
380 # In[40]:
381
382
383 # Confusion matrix
384 cm = confusion_matrix(y_test, y_pred, labels=pipeline.classes_)
385 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=pipeline.classes_)
386 fig, ax = plt.subplots(figsize=(15,7))
387 disp.plot(ax=ax)
388 plt.show()
389

```